

Gradient Descent as Implicit EM in Distance-Based Neural Models

Alan Oursland

December 2025

Abstract

Neural networks trained with standard objectives exhibit behaviors characteristic of probabilistic inference: soft clustering, prototype specialization, and Bayesian uncertainty tracking. These phenomena appear across architectures—in attention mechanisms, classification heads, and energy-based models—yet existing explanations rely on loose analogies to mixture models or post-hoc architectural interpretation. We provide a direct explanation. For any objective with log-sum-exp structure over distances or energies, the gradient with respect to each distance is exactly the negative posterior responsibility of the corresponding component: $\partial L / \partial d_j = -r_j$. This is an algebraic identity, not an approximation—a specialization of Fisher’s identity from the classical EM literature, whose significance here is its address: standard neural objectives instantiate it without modification. The immediate consequence is that gradient descent on such objectives performs expectation-maximization implicitly—responsibilities are not auxiliary variables to be computed but gradients to be applied. No explicit inference algorithm is required because inference is embedded in optimization. This result unifies three regimes of learning under a single mechanism: unsupervised mixture modeling, where responsibilities are fully latent; attention, where responsibilities are conditioned on queries; and cross-entropy classification, where supervision clamps responsibilities to targets. Our claims live at training time: the responsibility-weighted gradient dynamics recently documented in transformers are forced by the objective’s geometry. The in-context Bayesian computation that trained transformers perform at inference time is the endpoint of these dynamics, not their per-step content. For this class of objectives, learning and inference are the same process.

1 Introduction

Neural networks trained with standard objectives repeatedly exhibit behaviors associated with probabilistic inference: soft clustering, prototype specialization, uncertainty tracking, and mixture-model dynamics. These phenomena appear across architectures—in attention mechanisms, classification heads, and energy-based models—yet their origin has remained unclear. Are they emergent properties of scale? Architectural accidents? Or something more fundamental? This paper argues that such behaviors are none of these. They are necessary consequences of the geometry of common objective functions.

1.1 The Puzzle

Consider the range of phenomena that appear, unbidden, in trained neural networks. Attention heads in transformers learn to specialize, each routing information for distinct semantic roles. Classification networks partition representation space into regions that behave like mixture components. Deep networks trained on noisy data exhibit robustness patterns reminiscent of Bayesian inference, down-weighting outliers and tracking uncertainty across inputs. These behaviors emerge without explicit

probabilistic modeling, without mixture-model architectures, and without any algorithm resembling expectation-maximization.

The standard explanations are unsatisfying. One view holds that these are emergent properties of scale—that sufficient parameters and data somehow give rise to statistical structure. Another treats them as architectural coincidences, artifacts of specific design choices like softmax normalization or residual connections. A third offers loose analogies: attention is “like” soft clustering; cross-entropy “resembles” a mixture model. None of these explanations answers the deeper question: *why do these specific behaviors appear, rather than others?* And why do they appear so reliably, across such different architectures and tasks?

1.2 Recent Evidence

Recent work by Agarwal et al. [2025a,b] has sharpened this puzzle considerably. In controlled experimental settings—“Bayesian wind tunnels” where the true posterior is analytically known—small transformers match the analytic posterior’s entropy, position by position, to within 10^{-3} – 10^{-4} bits: agreement at the resolution of the arithmetic. Capacity-matched MLPs trained under identical conditions fail catastrophically, suggesting that the phenomenon depends on the inductive biases of attention rather than optimization alone.

More striking still is what Agarwal et al. find in the gradient dynamics. Attention weights stabilize early in training while value vectors continue to refine—a two-timescale structure that mirrors the E-step and M-step of classical expectation-maximization. Values receive updates weighted by attention, precisely as prototypes receive updates weighted by responsibilities in mixture models. The authors provide a complete first-order analysis showing that this structure is not incidental but systematic.

A scoping distinction is essential to everything that follows. Two different processes are in play. The *Bayesian inference* their wind tunnels demonstrate is in-context computation: a frozen, trained network tracking the posterior predictive over latent task variables as a sequence unfolds, at inference time—a point made explicit in a clarification added in revised versions of both papers. The *EM-like dynamics* are a training-time phenomenon: their gradient analysis, in their words, “lives at the EM/SGD level,” explaining how cross-entropy training sculpts the geometry that the frozen network later uses. The present paper’s claims live entirely at the second level—training time. We analyze what gradient descent does, not what the learned function computes in context.

Within that training-time level, Agarwal et al. are careful about the status of the EM correspondence. Their gradient laws are derived, not analogized—the mechanics are exact consequences of cross-entropy. What they decline is the probabilistic interpretation: because value updates are driven by the backpropagated error signal rather than by observed data, values move to explain the error geometry rather than to maximize a likelihood over inputs, and so, in their words, “the analogy is structural rather than variational.” What their account leaves open is the probabilistic status of the training dynamics: whether the responsibilities that gate learning are posteriors of anything.

1.3 This Paper

Before stating what this paper shows, it is worth stating where it came from, because the route explains the framing. The result was not found by interpreting a trained network. It was found by taking the distance-based reading of neural outputs [Oursland, 2024b] as a design constraint and asking the engineering question it forces: *if a network’s outputs are distances, what loss function should consume them?* Following the types—exponentiate a distance to get a likelihood, normalize

to compare alternatives, take the logarithm to return to the distance domain—leads directly to the log-sum-exp objective, the smooth “distance to the nearest prototype” (Section 2.2). Differentiating that objective to understand its training behavior is what exposed the identity below: the mixture model was not imposed on the loss; it was discovered inside it. The paper is organized to preserve that direction of travel: geometry first, probability derived.

This paper supplies the probabilistic status that Agarwal et al.’s caveat leaves open, and delimits it. We show that for objectives with log-sum-exp structure over distances or energies, the gradient of the loss with respect to each distance is exactly the posterior responsibility of the corresponding component. This is not an approximation, not a resemblance, not an analogy. It is an algebraic identity:

$$\frac{\partial L}{\partial d_j} = -r_j \tag{1}$$

The identity itself is classical—it is Fisher’s identity from the EM literature, specialized to log-sum-exp objectives—and the precise relationship between gradient descent and EM has been characterized before [Xu and Jordan, 1996, Neal and Hinton, 1998]. What has not been drawn out is where this identity lives: standard neural objectives instantiate it, unmodified, over the distance-based representations that ordinary networks compute. When the log-sum-exp is the loss, the EM correspondence is therefore not structural but variational: the responsibilities are posterior probabilities of an implicit mixture, and gradient descent performs expectation-maximization implicitly, continuously rather than in discrete alternating steps. The forward pass computes unnormalized likelihoods; normalization yields responsibilities; backpropagation delivers responsibility-weighted updates to parameters. No auxiliary latent variables are introduced. No inference algorithm is invoked. The gradient *is* the responsibility. When the softmax instead sits inside the network—attention—the same calculus recovers exactly the structural laws of Agarwal et al. [2025b], and explains why the variational reading applies at the loss level and becomes structural in the interior (Section 4.2).

This reframes the relationship between optimization and inference. Inference is not a separate algorithmic layer added on top of learning. It is not a post-hoc interpretation of learned representations. Under the objectives we analyze, inference and optimization are the same computation viewed at different levels of abstraction. The Bayesian structure observed by Agarwal et al. is not an emergent property that happens to appear; it is forced by the geometry of the loss.

1.4 Contributions

The contributions of this paper are deliberately narrow and can be summarized as three claims at increasing levels of abstraction.

One theorem. For any objective of the form $L = \log \sum_j \exp(-d_j)$, the gradient with respect to each distance satisfies $\partial L / \partial d_j = -r_j$, where r_j is the normalized exponential—the responsibility of component j . The identity is a specialization of Fisher’s identity from the classical EM literature (Section 5.2); it requires no conditions beyond differentiability. What is new is not the identity but its address: standard neural objectives—cross-entropy, attention, mixture likelihoods—instantiate it without modification.

One interpretation. This identity implies that gradient descent on distance-based log-sum-exp objectives performs implicit expectation-maximization. The E-step is the forward pass; the M-step is the parameter update; responsibilities are never explicitly computed because they *are* the gradients. EM is not approximated by neural training—it is realized by it.

One unification. The same mechanism manifests across three regimes depending on the constraints imposed. In the *unsupervised* regime, responsibilities are fully latent and prototypes compete freely. In the *conditional* regime—attention—responsibilities are recomputed per query over a shared prototype family. In the *constrained* regime—cross-entropy classification—supervision clamps one responsibility while competition persists among alternatives. These are not three different phenomena. They are one phenomenon under different boundary conditions.

2 The Geometric Substrate

The main result of this paper—that responsibilities emerge as gradients—rests on a specific geometric foundation. This section establishes that foundation. We first summarize the interpretation of neural network outputs as distances or energies rather than confidences, importing this result from prior work. We then define the class of log-sum-exp objectives to which our analysis applies. Finally, we briefly recall the structure of classical expectation-maximization, not because our method requires it, but to establish the reference point against which implicit EM will be compared.

2.1 Distance-Based Representations

The standard interpretation of neural network outputs treats them as confidences or scores indicating the strength of evidence for a hypothesis. A logit is high when the network “believes” in a class; an attention score is high when a query “matches” a key. This interpretation, while intuitive, obscures the geometric structure of what neural networks actually compute.

In prior work [Oursland, 2024b], we developed a different interpretation: the outputs of standard neural networks are better understood as distances or energies relative to learned prototypes. A linear layer followed by an absolute value or ReLU activation computes a quantity that behaves as a (possibly signed) deviation from a learned reference—a Mahalanobis distance along a principal direction. Concretely, the Mahalanobis distance of a point x from a Gaussian component with mean μ and principal direction v scaled by eigenvalue λ is $|\lambda^{-1/2}v^\top(x - \mu)|$, which has the form $|Wx + b|$ for appropriate W and b . Standard ReLU networks compute this via the identity $|z| = \text{relu}(z) + \text{relu}(-z)$, decomposing signed distance into two half-space detectors. Under this view, a low output indicates proximity to a prototype; a high output indicates deviation. Probabilities are not primitive quantities but derived ones, arising only after exponentiation and normalization transform distances into relative likelihoods. Empirical companions to that work test the distance-metric hypothesis directly [Oursland, 2024a, 2025].

This interpretation reflects the structure of the computation itself: the weights of a linear layer define a basis; the biases define offsets along that basis; the activation measures deviation. What changes is not the computation but the semantics we assign to it. Throughout this paper, we adopt the distance-based interpretation and refer to neural outputs as energies or distances interchangeably. We emphasize, however, that the main result does not depend on this reading: the gradient identity of Section 3 is purely algebraic and holds for arbitrary energies. The distance interpretation supplies the semantics under which that identity deserves the name *inference*—it motivates the framework rather than propping up the mathematics.

2.2 The Log-Sum-Exp Objective

Given a set of distances or energies $\{d_1, d_2, \dots, d_K\}$ computed for an input x , we consider objectives of the form:

$$L(x) = \log \sum_{j=1}^K \exp(-d_j(x)) \quad (2)$$

This objective admits two readings, and the order in which they are taken matters to this paper.

The first reading is geometric and requires no probability at all. The negative of L is the *soft minimum* of the distances: for K components,

$$\min_j d_j - \log K \leq -L \leq \min_j d_j$$

so minimizing $-L$ minimizes, up to an additive constant, the distance from the input to the *nearest* prototype. The LSE objective is the smooth answer to a design question: given a set of distances, what loss says “be close to at least one of them”? It is the distance-to-a-set, softened enough to be differentiable. Under this reading, the gradient concentration of Theorem 1—signal flowing to the closest component—is simply the smooth analogue of the fact that a minimum is attained at its argmin. This design-first route is, in fact, how the present framework was arrived at: starting from distance-based representations and asking what objective should consume a distance, one is led to exponentiation (distance to likelihood), normalization (comparison across alternatives), and the logarithm (back to the distance domain)—at which point the mixture model is discovered already inside the loss, not imposed on it.

The second reading is probabilistic: if $\exp(-d_j)$ represents the unnormalized likelihood that component j generated the input, then L is the log marginal likelihood—the log-probability that *some* component generated it. Maximizing L encourages the model to place at least one prototype close to each input. The two readings pick out the same objective; the geometric one explains why it is natural, the probabilistic one explains what its gradients mean.

We adopt the following notation throughout. Let $P_j = \exp(-d_j)$ denote the unnormalized likelihood of component j . Let $Z = \sum_k P_k = \sum_k \exp(-d_k)$ denote the partition function. Define the responsibility of component j as:

$$r_j = \frac{P_j}{Z} = \frac{\exp(-d_j)}{\sum_k \exp(-d_k)} \quad (3)$$

The responsibilities are non-negative and sum to one. They represent the posterior probability that component j is responsible for the input, under a uniform prior over components.

This structure appears throughout deep learning under various names. In cross-entropy classification, the logits z_j act as negative distances ($d_j = -z_j$), and the objective becomes $L = -z_y + \log \sum_k \exp(z_k)$, where y is the label. In attention mechanisms, the scores s_{ij} act as negative distances between query i and key j , and softmax normalization produces responsibilities that weight the value vectors. These are not analogies to the LSE objective; they are instances of it.

2.3 Classical EM

Expectation-maximization is the classical algorithm for fitting mixture models with latent assignments. It proceeds in two alternating steps.

In the E-step, responsibilities are computed. Given current parameters, each data point is softly assigned to each component based on relative likelihood:

$$r_j^{(t)} = \frac{P(x | \theta_j)}{\sum_k P(x | \theta_k)}$$

These responsibilities sum to one and represent the posterior probability that component j generated the observation.

In the M-step, parameters are updated. Each component’s parameters are adjusted to better fit the data points assigned to it, weighted by responsibility:

$$\theta_j^{(t+1)} = \arg \max_{\theta_j} \sum_i r_{ij}^{(t)} \log P(x_i | \theta_j)$$

For Gaussian mixtures, this reduces to computing responsibility-weighted means and covariances. The key property is that every data point influences every component, but the influence is gated by how much responsibility that component holds for that point.

Classical EM is discrete and alternating: compute all responsibilities, then update all parameters, then repeat. The E-step and M-step are separate procedures with distinct computational roles. This separation is algorithmic, not fundamental: [Neal and Hinton \[1998\]](#) showed that both steps are coordinate ascent on a single free-energy objective, a view whose consequences for gradient-based training we develop in what follows.

3 Main Result: Responsibilities Are Gradients

We now state and derive the central result. The derivation is elementary—a single application of the chain rule—but its implications are significant. We show that for log-sum-exp objectives over distances, the gradient with respect to each distance is exactly the negative responsibility of the corresponding component. This identity requires no approximations and holds for any model that computes distances and optimizes an LSE objective via gradient descent. The identity itself is classical: it is Fisher’s identity—the gradient of an incomplete-data log-likelihood equals the posterior-expected complete-data gradient [[Cappé and Moulines, 2009](#), [Bishop, 2006](#)—specialized to a uniform-prior mixture over exponentiated distances (Section 5.2). The contribution of this paper is not the derivative but the recognition of where it applies: standard neural objectives instantiate this structure over distance-based representations without any modification, so every gradient step of ordinary training is a responsibility-weighted update.

3.1 Derivation

Let $d_j(x)$ denote the distance or energy associated with component j for input x . Define unnormalized likelihoods $P_j = \exp(-d_j)$ and the partition function $Z = \sum_k \exp(-d_k)$. Consider the log-sum-exp objective:

$$L = \log \sum_{j=1}^K \exp(-d_j) = \log Z \tag{4}$$

We compute the gradient of L with respect to a single distance d_i :

$$\frac{\partial L}{\partial d_i} = \frac{1}{Z} \cdot \frac{\partial}{\partial d_i} \sum_{j=1}^K \exp(-d_j)$$

Only the i -th term in the sum depends on d_i :

$$\frac{\partial L}{\partial d_i} = \frac{1}{Z} \cdot \frac{\partial}{\partial d_i} \exp(-d_i) = \frac{1}{Z} \cdot (-1) \cdot \exp(-d_i)$$

Simplifying:

$$\frac{\partial L}{\partial d_i} = -\frac{\exp(-d_i)}{\sum_k \exp(-d_k)}$$

The right-hand side is the negative of the normalized exponential—precisely the responsibility r_i as defined in Section 2.2. This yields the main result.

Theorem 1. *For any objective of the form $L = \log \sum_j \exp(-d_j)$, the gradient with respect to the j -th distance is the negative responsibility of component j :*

$$\frac{\partial L}{\partial d_j} = -r_j \quad \text{where} \quad r_j = \frac{\exp(-d_j)}{\sum_k \exp(-d_k)} \quad (5)$$

No approximations have been made. The identity holds exactly whenever the objective has LSE form and the distances are differentiable.

3.2 What This Means

The theorem admits a direct interpretation: responsibilities are not quantities that must be computed and stored. They are gradients. Any system that computes distances, applies a log-sum-exp objective, and updates parameters via gradient descent is already computing responsibilities—not as an intermediate step, but as the learning signal itself.

Note that the identity itself is purely algebraic—it holds for any log-sum-exp objective regardless of how one interprets the quantities involved. The EM interpretation, however, depends on reading the d_j as distances and the r_j as posterior responsibilities. The mathematics forces the gradient structure; the semantics give it meaning.

Consider what happens during training. The forward pass computes distances d_j and evaluates the objective $L = \log \sum_j \exp(-d_j)$. This computation implicitly determines the responsibilities $r_j = \exp(-d_j)/Z$, but they are never explicitly represented. Backpropagation then computes $\partial L / \partial d_j = -r_j$, and the chain rule propagates this gradient to the parameters that produced d_j . The parameters of component j receive gradient signal proportional to r_j —the responsibility of that component for the input.

This is exactly the M-step of expectation-maximization. In classical EM, responsibilities are computed in the E-step and then used to weight parameter updates in the M-step. Here, no E-step is required as a separate operation. The forward pass *is* the E-step; the backward pass *is* the M-step. The two are fused into a single gradient computation.

The consequence is that gradient descent on LSE objectives performs EM implicitly. The precise relationship is inherited from classical results: the gradient is the responsibility-weighted update direction (Fisher’s identity), and the EM step is that same gradient under a positive-definite preconditioner [Xu and Jordan, 1996]—the two methods share update structure, though not trajectories [Zhang et al., 2020]. The discrete alternation of classical EM—compute responsibilities, then update parameters, then repeat—is likewise not fundamental: E and M are coordinate ascent on a single objective [Neal and Hinton, 1998], and under gradient descent they collapse into continuous, simultaneous optimization. Every gradient step is a responsibility-weighted update. Every trained network has been doing responsibility-weighted learning all along.

Throughout this paper, “implicit EM” refers to the emergence of responsibility-weighted parameter updates under gradient descent on log-sum-exp objectives—not to coordinate-ascent EM or guarantees about convergence.

3.3 Conditions

The result depends on three structural conditions. When all three are satisfied, implicit EM dynamics are not optional—they are forced by the mathematics.

Exponentiation. The distances must be passed through an exponential to form unnormalized likelihoods $P_j = \exp(-d_j)$. This transformation converts additive differences in distance into multiplicative ratios in likelihood. Without exponentiation, the gradient with respect to d_j does not yield a normalized quantity, and responsibilities do not emerge.

Normalization. The likelihoods must be normalized across alternatives, either explicitly through softmax or implicitly through the log-sum-exp structure. Normalization induces competition: increasing the likelihood of one component necessarily decreases the relative likelihood of others. This competition is what makes responsibilities sum to one and what causes gradient signal to be distributed according to relative fit. Without normalization, components operate independently, and there is no assignment structure.

Gradient-based optimization. The parameters must be updated via gradients of the objective. The identity $\partial L / \partial d_j = -r_j$ is a fact about the loss surface; it becomes a fact about learning dynamics only when gradient descent (or a variant) is used. Other optimization methods—evolutionary strategies, random search, discrete updates—would not automatically inherit the responsibility-weighted structure.

When these conditions hold, there is no additional design choice that enables or disables EM-like behavior. It is a consequence of the objective geometry. Any architecture that computes distances, normalizes via exponentiation, and trains with gradients will exhibit implicit EM.

3.4 Completing the Gaussian: The Volume Term

The analysis so far treats the distances d_j as given, which corresponds to mixture components of fixed shape. A full Gaussian likelihood carries an additional normalization term—the volume of each component’s covariance—and neural objectives that omit it are exposed to metric collapse (Section 6.2). For the common case where distances are computed by linear layers, this missing term is itself expressible in network primitives, and Theorem 1 extends to it verbatim.

Let component j compute $z_j = W_j x + b_j$ with W_j square and full-rank, and define the *volume-corrected energy*:

$$\tilde{e}_j = \frac{1}{2} \|z_j\|^2 - \log |\det W_j| \quad (6)$$

Then $\exp(-\tilde{e}_j) = |\det W_j| \exp(-\frac{1}{2} \|W_j x + b_j\|^2)$, which is, up to the constant $(2\pi)^{-D/2}$ shared by all components, exactly the density $\mathcal{N}(x; \mu_j, \Sigma_j)$ with $\mu_j = -W_j^{-1} b_j$ and $\Sigma_j^{-1} = W_j^\top W_j$. The identity $|\Sigma_j|^{-1/2} = |\det W_j|$ is the change-of-variables Jacobian familiar from normalizing flows [Rezende and Mohamed, 2015]. The log-sum-exp objective over volume-corrected energies, $L = \log \sum_j \exp(-\tilde{e}_j)$, is therefore the exact log marginal likelihood of a full Gaussian mixture—no proportionality, no fixed-covariance assumption—and Theorem 1 applies unchanged: $\partial L / \partial \tilde{e}_j = -r_j$, where r_j is now the exact posterior responsibility under fully normalized components.

Two consequences follow. First, implicit EM extends from means to covariances. The gradient of the objective with respect to the weight matrix is, per sample,

$$\frac{\partial L}{\partial W_j} = r_j (W_j^{-\top} - z_j x^\top) \quad (7)$$

which is the responsibility times the gradient of the complete-data Gaussian log-likelihood—Fisher’s identity at the parameter level. Gradient descent on this objective performs responsibility-

weighted updates to the full component geometry, exactly as the M-step of full-covariance EM does.

Second, the volume term removes the trivial collapse mode. Without it, the objective $\exp(-d_j)$ alone is maximized by shrinking the learned metric—driving $W_j \rightarrow 0$ maps every input to the prototype, sending all distances to zero and all likelihoods to one. The $\log|\det W_j|$ term sends that configuration to $-\infty$, and equilibrium is reached where the learned metric matches the data’s covariance: growing W_j is rewarded by volume but punished quadratically by distance, and vice versa. The classical single-point singularity of Gaussian mixture MLE—a component locking onto one data point and shrinking its variance without bound—remains, as it does in classical EM [Bishop, 2006], and is mitigated in practice by batch averaging or bounded parameterizations. Computationally, the determinant is cheap under standard parameterizations: for diagonal or triangular W_j , $\log|\det W_j|$ is a sum of log-diagonals, at $O(D)$ cost.

The point is not that practitioners should always add this term. It is that the framework is constructive as well as diagnostic: the same primitives that compute distances can compute the volume correction, the corrected objective is the exact mixture likelihood, and the implicit EM identity survives intact. Where Section 6.2 explains why omitting the term invites collapse, this section shows what including it costs—one log-determinant—and what it buys.

4 Three Regimes of Implicit Inference

The same mechanism manifests differently under different constraints:

4.1 Unsupervised Regime: Mixture Learning

In the purest case, the log-sum-exp objective operates without supervision. A model computes distances $d_j(x)$ from an input to each of K components, and training minimizes:

$$L = -\log \sum_{j=1}^K \exp(-d_j) \quad (8)$$

This is the negative log marginal likelihood—the objective used in classical mixture model fitting. All components compete for every input. No labels constrain which component should win.

By Theorem 1, the gradient of this minimization objective with respect to each distance is the responsibility itself: $\partial L / \partial d_j = +r_j$. (The sign flip from Section 3 reflects the switch from maximizing log-likelihood to minimizing its negative; the dynamics are identical.) A gradient step therefore decreases each distance in proportion to its responsibility. Components that are close to an input (low d_j , high r_j) receive strong gradient signal to move closer. Components that are far (high d_j , low r_j) receive weak signal and remain largely unchanged. This is precisely the M-step dynamic of Gaussian mixture models: prototypes update toward data points in proportion to their responsibility for those points.

The result is spontaneous specialization. Even with random initialization, components differentiate over training. Each prototype drifts toward a region of input space for which it consistently takes high responsibility, while ceding other regions to competitors. Clustering emerges not because it was specified but because the objective geometry forces responsibility-weighted updates.

This regime corresponds exactly to classical EM on mixture models. The fixed points—and the path toward them—are governed by the same responsibilities.

4.2 Conditional Regime: Attention Mechanisms

Attention mechanisms instantiate the same structure in a conditional setting. For a query q_i and a set of keys $\{k_j\}$, attention scores are computed as:

$$s_{ij} = \frac{q_i^\top k_j}{\sqrt{d_k}}$$

These scores act as negative distances: high scores indicate proximity in the query-key space. Softmax normalization converts scores to attention weights:

$$\alpha_{ij} = \frac{\exp(s_{ij})}{\sum_m \exp(s_{im})}$$

The attention weights satisfy the definition of responsibilities exactly. They are non-negative, sum to one across keys, and represent the degree to which each key “explains” the query. The output is a responsibility-weighted combination of values: $o_i = \sum_j \alpha_{ij} v_j$. This is not an analogy to mixture modeling. It is mixture modeling, conditioned on the query.

An honest distinction is required here, and it sharpens the framework rather than weakening it. Theorem 1 applies when the log-sum-exp is the *objective*. In a transformer, the attention softmax sits *inside the network*; the loss is cross-entropy at the output. The gradient of the loss with respect to an attention score is therefore not $-\alpha_{ij}$. The exact identities are obtained by the chain rule. Let $g_i = \partial L / \partial o_i$ be the error backpropagated to the attention output, and define the alignment $a_{ij} = g_i^\top v_j$ with attention-weighted mean $\bar{a}_i = \sum_k \alpha_{ik} a_{ik}$. Then:

$$\frac{\partial L}{\partial v_j} = \sum_i \alpha_{ij} g_i \quad \text{and} \quad \frac{\partial L}{\partial s_{ij}} = \alpha_{ij} (a_{ij} - \bar{a}_i) \quad (9)$$

The first identity is exact and has the same strength as Theorem 1: each value receives the downstream error scaled precisely by its responsibility. This is the M-step for internal softmaxes, holding verbatim at every gradient step for any loss. The second identity shows that scores follow an advantage rule: a score rises only if its value is more useful to the downstream error than the current attention-weighted average. Agarwal et al. [2025b] derive exactly these two laws for transformer attention; they hold for any softmax module. Responsibilities still gate all learning—both gradients carry the factor α_{ij} —but where the loss-level LSE competes on *proximity*, the internal LSE competes on *usefulness*. Two consequences follow immediately. Score gradients sum to zero across keys: internal competition redistributes and cannot uniformly favor all components. And once attention concentrates ($\alpha_{ij^*} \rightarrow 1$), the advantage $a_{ij^*} - \bar{a}_i \rightarrow 0$, so score updates vanish while value updates persist—the two-timescale dynamic observed empirically by Agarwal et al. is a one-line consequence of the gradient forms.

There is also a second, exact sense in which attention computes an LSE gradient—in the forward pass itself. The gradient of $\log \sum_j \exp(q_i^\top k_j)$ with respect to the query is $\sum_j \alpha_{ij} k_j$: the attention read-out is the gradient of a log-sum-exp energy, the retrieval dynamics of modern Hopfield networks [Ramsauer et al., 2021]. The forward pass takes an LSE gradient with respect to the *state* (inference); the backward pass takes responsibility-gated gradients with respect to the *parameters* (learning).

A second clarification concerns the prototypes. In a Gaussian mixture model, the component means are persistent parameters. In attention, the value vectors $v_j = W_V x_j$ are input-derived and transient. The persistent object is the value projection: by the chain rule, $\partial L / \partial W_V = \sum_{i,j} \alpha_{ij} g_i x_j^\top$ —the projection accumulates exactly responsibility-weighted outer products. This is the M-step applied not to fixed prototypes but to the parameters of a prototype-generating function. The EM dynamics thus operate across training, not within a single forward pass. Over the course of training,

W_V specializes: it learns to produce values that are useful precisely to the queries that attend to them.

Attention is conditional mixture inference. The query specifies the context; the keys define candidate components; the values represent what each component contributes. The mixture is re-instantiated for every query, but the underlying parameters— W_Q , W_K , W_V —accumulate responsibility-weighted updates across all queries and all training examples. This is implicit EM in function space rather than data space.

4.3 Constrained Regime: Cross-Entropy Classification

Cross-entropy classification introduces supervision, but it does not escape the implicit EM framework. It constrains it.

In standard classification, the model produces logits z_j for each class. Interpreting logits as negative distances ($d_j = -z_j$), the cross-entropy loss for a sample with label y is:

$$L = -z_y + \log \sum_k \exp(z_k) = d_y + \log \sum_k \exp(-d_k)$$

This is the log-sum-exp objective plus a term that penalizes the distance to the correct class. Computing the gradient with respect to d_j :

$$\frac{\partial L}{\partial d_j} = \mathbb{I}[j = y] - \frac{\exp(-d_j)}{\sum_k \exp(-d_k)} = \mathbb{I}[j = y] - r_j$$

(In logit form this is the familiar $\partial L / \partial z_j = p_j - \mathbb{I}[j = y]$; the sign flips because $d_j = -z_j$.) For incorrect classes ($j \neq y$), the gradient is $-r_j$: a descent step *increases* their distance in proportion to their current responsibility—they are pushed away exactly as much as they presume to explain the input. For the correct class ($j = y$), the gradient is $1 - r_y \geq 0$: a descent step decreases its distance, with strength proportional to how much responsibility it is *missing*.

The label acts as a clamp. In unsupervised mixture learning, responsibilities are entirely latent—the data determines which component wins. In cross-entropy, the label declares that the correct class *should* have responsibility 1. The gradient then corrects any deviation from this target. If the correct class already dominates ($r_y \approx 1$), the gradient is small. If incorrect classes have stolen probability mass ($r_y < 1$), the gradient is large.

Competition among incorrect classes remains intact. When the model misclassifies, the responsibility mass is distributed among wrong answers, and each receives gradient signal proportional to its share. The classes that are “most wrong”—those with highest r_j —are penalized most strongly. This is not uniform repulsion; it is responsibility-weighted correction.

Cross-entropy does not eliminate the EM dynamics; it directs them. The M-step still updates components in proportion to their responsibilities, but supervision biases the process toward a predetermined assignment. This explains why cross-entropy is so effective despite its simplicity: it inherits the soft competition and automatic weighting of mixture models while channeling those dynamics toward a supervised target. The loss function is doing more work than its familiar form suggests.

4.4 The Taxonomy

The three regimes—unsupervised, conditional, constrained—differ in what is observed and what is latent. They also differ in where the log-sum-exp sits. When the LSE is the loss (mixture likelihood, cross-entropy), gradients *equal* responsibilities, signed by supervision. When the LSE

is inside the network (attention), value gradients are exactly responsibility-weighted while score gradients are responsibility-gated advantages. In every case the regimes share a common structure: exponentiation of distances followed by normalization across alternatives. This structure is what produces responsibilities, and responsibilities are what produce implicit EM—what varies is which quantity the responsibilities weight.

The critical factor is normalization. When outputs are normalized—whether by softmax, by the log-sum-exp partition function, or by any operation that forces a sum-to-one constraint—components compete. An increase in one component’s likelihood necessarily decreases the relative likelihood of others. This competition is the source of assignment: each input is probabilistically allocated across components, and gradients distribute accordingly.

Remove normalization, and the structure collapses. Consider objectives based on Gaussian kernels without a partition function, such as maximum correntropy:

$$L = 1 - \exp(-d^2/\sigma^2)$$

Here, each component operates independently. A point far from all prototypes produces vanishing gradients for all of them—not because responsibility has been assigned elsewhere, but because no competition exists to distribute. There is no implicit E-step because there are no responsibilities. The objective gains robustness to outliers (points far from all prototypes are effectively ignored) but loses the assignment structure entirely.

This clarifies the design space. Exponentiation converts distances to likelihoods; normalization converts likelihoods to responsibilities. With both, implicit EM is unavoidable. With exponentiation alone, the model gains robustness but abandons inference. The choice of objective is a choice about whether the model should assign or ignore—and that choice is made at the level of the loss function, not the architecture.

5 Relation to Prior Work

The result derived here does not stand alone. It depends on prior work for its geometric foundation and gains significance from recent empirical findings that it explains. This section positions the contribution relative to four bodies of work: the distance-based interpretation of neural representations, classical results connecting EM to gradient methods, recent evidence of Bayesian structure in transformers, and energy-based learning.

5.1 Prior Work on Distance-Based Representations

In prior work [Oursland, 2024b], we established the geometric interpretation on which this paper depends. That work shows that standard neural network layers—affine transformations followed by ReLU or absolute value activations—compute quantities that behave as distances from learned prototypes. Outputs are deviations, not confidences; they measure how far an input lies from a reference structure defined by the weights. This interpretation is not imposed but derived from the mathematics of the operations involved.

The present work takes this geometric substrate as given and asks a different question: what happens when distance-based representations are optimized under standard objectives? Our prior work is silent on learning dynamics. It characterizes what neural networks represent, not how they learn. The contribution here is to show that log-sum-exp objectives over distances produce responsibility-weighted gradients, and that this induces implicit EM.

The two results are complementary and together form a complete picture. The first says: neural networks compute distances. The second says: optimizing distances with LSE objectives performs inference. Neither subsumes the other. Without the distance interpretation, the identification of gradients with responsibilities is a formal curiosity. Without the gradient identity, the distance interpretation describes static representations with no account of how they emerge. The geometric substrate enables the learning dynamics; the learning dynamics justify attention to the geometric substrate.

5.2 Classical Results Connecting EM and Gradient Methods

The identity at the center of this paper is not new to statistics. It is a specialization of *Fisher’s identity*: for any latent-variable model, the gradient of the incomplete-data log-likelihood equals the posterior expectation of the complete-data gradient [see, e.g., [Cappé and Moulines, 2009](#)]. For a mixture model, that posterior is exactly the vector of responsibilities, so marginal-likelihood gradients are responsibility-weighted by construction; the Gaussian-mixture case is textbook material [[Bishop, 2006](#), §9.2]. Theorem 1 is Fisher’s identity for a uniform-prior mixture with component likelihoods $\exp(-d_j)$, differentiated with respect to the distances themselves.

The relationship between gradient descent and EM is also precisely characterized. [Xu and Jordan \[1996\]](#) prove that the EM update for Gaussian mixtures is a preconditioned gradient step on the log-likelihood, with an explicit positive-definite projection matrix; [Salakhutdinov et al. \[2003\]](#) use the expected-gradient identity to build direct optimizers for latent-variable models and characterize when EM behaves like a quasi-Newton method. [Neal and Hinton \[1998\]](#) show that the E-step and M-step are coordinate ascent on a single free-energy objective, so the discrete alternation of classical EM is an algorithmic choice rather than a structural fact. Our claim that inference and optimization fuse under LSE objectives inherits its precise form from these results: the instantaneous gradient is the responsibility-weighted update direction (Fisher), and the EM step is that same gradient under a positive-definite rescaling (Xu & Jordan).

At the architectural level, the EM/attention connection also has a lineage. Modern Hopfield networks [[Ramsauer et al., 2021](#)] identify transformer attention with retrieval dynamics on a log-sum-exp energy—the attention output is the gradient of an LSE with respect to the query. Neural Expectation Maximization [[Greff et al., 2017](#)], EM attention modules [[Li et al., 2019](#)], and EM capsule routing [[Hinton et al., 2018](#)] all build explicit EM iterations into networks.

Given this literature, the contribution of the present paper is deliberately not the identity itself. It is the recognition that *standard* neural objectives—softmax cross-entropy, attention, mixture likelihoods—already instantiate Fisher’s identity over distance-based representations without modification, and the analysis of how one mechanism manifests across the unsupervised, conditional, and constrained regimes. The classical results concern generative latent-variable models optimized deliberately; we argue the same structure operates, unnoticed, inside ordinary discriminative training.

5.3 Agarwal et al. (2025a,b)

[Agarwal et al. \[2025a\]](#) provide striking empirical evidence that transformers implement Bayesian inference. In controlled settings where the true posterior is analytically known—their “Bayesian wind tunnels”—small transformers match the analytic posterior’s entropy, position by position, to within 10^{-3} – 10^{-4} bits, while capacity-matched MLPs fail catastrophically. They also prove the endpoint statement: the population optimum of cross-entropy is the Bayes posterior predictive. That theorem is architecture-agnostic—it identifies what the optimum is, not how gradient training reaches it or what the intermediate quantities mean along the way.

Agarwal et al. [2025b] move from statics to dynamics, analyzing the gradients of attention under cross-entropy training. They derive an advantage-based routing law for attention scores and a responsibility-weighted update for value vectors—the same identities we state in Section 4.2. They observe a two-timescale structure: attention patterns stabilize early while values continue to refine—mirroring the E-step and M-step of classical EM. The analysis is thorough and the EM parallel is explicit.

The authors are careful, however, about the status of the EM correspondence. Their gradient laws are derived, not analogized—the mechanics are exact consequences of cross-entropy. What they decline is the probabilistic interpretation. Because the value updates are driven by the backpropagated error signal rather than by observed data, they note that values move to better explain the error geometry “rather than to maximize a likelihood over inputs. The analogy is structural rather than variational.” In their framing, the EM connection is a mechanistic correspondence, not the optimization of any latent-variable likelihood.

The present work supplies the probabilistic status that this caveat leaves open—and delimits it. When the log-sum-exp is the loss (a mixture likelihood, or cross-entropy at the output layer), the correspondence is genuinely variational: the gradient with respect to each distance is the posterior responsibility of an implicit mixture (Fisher’s identity), and implicit EM is an identity, not a resemblance. When the softmax sits inside the network, Section 4.2 recovers exactly their two laws and gives their caveat precise form: internal competition is over downstream usefulness rather than likelihood, which is why the variational reading stops at the loss level. Their population-optimum theorem is the endpoint statement of the loss-level case; the responsibility-gradient identity is its per-step counterpart.

A boundary must also be drawn. In later revisions, the authors clarify that the Bayesian inference their experiments demonstrate is in-context computation at inference time—the posterior predictive over latent task variables, tracked by frozen weights as a sequence unfolds—and a third paper extends the geometric findings to production language models [Agarwal et al., 2025c]. Our framework does not explain that in-context computation. The responsibility-gradient identity is a statement about parameter updates during training; it connects to the inference-time geometry only through the endpoint theorem—training under cross-entropy drives the network toward the Bayes posterior predictive—and through their Paper II’s demonstration that the training dynamics build the geometric machinery. Explaining how a frozen transformer performs sequential belief updating in context is their contribution, not ours.

The relationship between the contributions is therefore complementary and level-separated. Agarwal et al. establish the empirical geometry, the in-context inference it supports, and the exact mechanics of the training gradients. We establish what those training gradients are probabilistically, and the conditions under which the EM reading is an identity rather than a correspondence.

5.4 Other Connections

The energy-based learning framework of LeCun et al. [2006] provides essential conceptual scaffolding. That work reframes learning as the minimization of an energy function, with probabilities derived via exponentiation and normalization. The log-sum-exp objective analyzed here is a special case of the “free energy” formulation in energy-based models, and the responsibility-gradient connection is the corresponding instance of Fisher’s identity discussed above. What we take from the energy-based view is the ordering of primitives: energies first, probabilities derived—the same ordering our distance-based reading assigns to neural outputs.

Dempster et al. [1977] introduced expectation-maximization as an algorithm for maximum likelihood estimation with latent variables. The E-step and M-step are defined as discrete, alternating

operations. The present work shows that for distance-based objectives, these steps collapse into gradient descent: the forward pass computes responsibilities implicitly, and backpropagation applies them. This does not contradict the classical formulation but reveals it as a special case of a more general phenomenon. EM is not merely an algorithm one can choose to apply; it is a property of certain objective geometries under gradient-based optimization.

Vaswani et al. [2017] introduced the transformer architecture with attention as its core mechanism. The original presentation emphasizes attention as a soft retrieval operation—queries attending to keys to aggregate values. The implicit EM perspective reframes attention as conditional mixture inference, with attention weights as responsibilities and value projections as prototype parameters. This interpretation is consistent with the original formulation but provides a probabilistic semantics that the architectural description lacks.

Mixture of Experts models [Jacobs et al., 1991] use explicit gating networks to route inputs to specialized subnetworks. The gating weights are responsibilities by another name. The difference is architectural: in Mixture of Experts, gating is a separate learned function; in standard attention and classification, responsibilities emerge as gradients of the objective without dedicated gating machinery. Implicit EM reveals that the explicit gating of Mixture of Experts is not necessary—any log-sum-exp objective produces responsibility-weighted routing automatically.

6 Limits and Failure Modes

The implicit EM framework is not universal. It applies under specific conditions, and when those conditions fail, so does the analysis. This section delineates the boundaries: when implicit EM does not arise, what pathologies can occur even when it does, and what phenomena lie entirely outside its scope. A clear account of limits strengthens rather than weakens the contribution—it distinguishes precise claims from overreach.

6.1 When Implicit EM Does Not Arise

The identity $\partial L / \partial d_j = -r_j$ depends on the log-sum-exp structure. Remove normalization, and the identity does not hold.

Consider multi-label classification with independent sigmoids. Each class j has its own output $\sigma(z_j)$, and the loss decomposes as a sum of independent binary cross-entropies. There is no partition function, no softmax, no competition across classes. The gradient with respect to z_j depends only on the target for class j , not on the outputs of other classes.

In this setting, responsibilities do not exist. No quantity sums to one across classes; no soft assignment distributes an input among alternatives. Each output channel operates in isolation. A point can be equally close to all prototypes or far from all of them, and the gradients do not redistribute—they simply reflect independent errors.

This is not a failure of the architecture but an absence of the required objective structure. Implicit EM arises from competition, and competition arises from normalization. Systems with independent outputs may learn useful representations, but they do not perform mixture inference and do not exhibit responsibility-weighted dynamics. The explanatory scope of the framework ends where normalization ends.

6.2 Scale and Collapse

A full Gaussian mixture model includes a log-determinant term in the likelihood—a penalty on the volume of each component’s covariance. This term prevents collapse: without it, a component could

shrink its covariance to zero, placing infinite density on a single point and achieving unbounded likelihood. The log-determinant diverges as the covariance collapses, counterbalancing the density gain.

Most neural objectives omit this term. Cross-entropy and attention softmax operate on distances or scores without explicit volume penalties. The implicit EM dynamics still hold—gradients are still responsibility-weighted—but nothing prevents the learned metric from degenerating. A network can learn to map all inputs to nearby points, collapsing the distance structure and trivializing the responsibilities. (Section 3.4 shows that for linear-layer distances the exact volume term is available in network primitives at the cost of one log-determinant; the analysis here concerns the common case where it is omitted.)

In practice, collapse is often avoided through implicit mechanisms: weight decay regularizes the scale of projections; layer normalization constrains activation magnitudes; architectural choices like residual connections preserve signal diversity. These interventions work, but they are not derived from the objective—they are heuristics that happen to stabilize geometry.

The implicit EM framework clarifies why collapse is a risk. When components are updated in proportion to their responsibilities, a component that captures slightly more mass receives stronger gradients, captures more mass still, and can dominate entirely. This positive feedback is inherent to EM dynamics and is classically controlled by the volume term. Neural networks remove that control and rely on other mechanisms to fill the gap. The framework does not solve this problem; it explains why the problem exists.

6.3 Supervision Constraints

In the unsupervised regime, responsibilities are entirely latent—the data alone determines which components claim which inputs. Supervision changes this. Labels declare which component *should* take responsibility, overriding what the geometry would otherwise dictate.

This constraint is powerful but rigid. Cross-entropy training forces the correct class toward responsibility 1, regardless of whether the input lies near that class’s prototype or far from all prototypes. An input equidistant from all class boundaries still receives a hard label; the model must assign it somewhere. The soft, graded structure of responsibilities persists among incorrect classes, but the correct class is clamped.

A consequence is the closed-world assumption. Softmax normalization guarantees that responsibilities sum to one—some class must take full responsibility for every input. There is no “none of the above,” no option for the model to reject an input as foreign to all known categories. An out-of-distribution input, no matter how anomalous, will be assigned to whichever class has the lowest distance, and the model’s confidence may be arbitrarily high.

This is a limitation of the objective, not the mechanism. The implicit EM framework describes how responsibilities arise and how they weight learning. It does not claim that the resulting assignments are semantically correct or that the objective captures all desirable behaviors. Cross-entropy with softmax enforces assignment; if non-assignment is needed, the objective must change. Open-set recognition, out-of-distribution detection, and selective prediction require objectives that break the closed-world assumption—either by introducing an explicit reject option or by abandoning normalization entirely. The implicit EM analysis applies to whatever objective is chosen; it does not choose the objective.

6.4 What This Framework Does Not Explain

The implicit EM framework explains one phenomenon: the emergence of responsibility-weighted learning dynamics from distance-based objectives. It does not explain everything neural networks do.

In-context, inference-time Bayesian computation is not addressed. A trained transformer with frozen weights can track a posterior over latent task variables as a sequence unfolds [Agarwal et al., 2025a,c]; that is a property of the learned function, executed in activation space at inference time. Our identity concerns gradient updates to parameters during training. The two are connected—training under cross-entropy drives the network toward the Bayes posterior predictive, and the training dynamics build the geometry the frozen network uses—but the framework derives the per-step structure of learning, not the algorithm the learned function implements in context. The forward-pass log-sum-exp gradient of Section 4.2 describes a single attention read-out, not sequential belief updating across a context.

Generalization—why networks perform well on unseen data—is not addressed. The framework describes the dynamics of training, not the inductive biases that enable generalization beyond the training distribution. Scaling laws—the predictable relationship between model size, data, and performance—lie outside the analysis entirely. Nothing in the gradient identity suggests how performance should scale with parameters or compute.

Long-horizon reasoning, planning, and sequential decision-making involve temporal structure that the framework does not capture. Implicit EM describes how a single input is softly assigned to components and how those components update. It does not describe how representations are composed over time, how goals propagate backward through action sequences, or how models learn to search.

Emergent capabilities—the sudden appearance of qualitatively new behaviors at scale—remain unexplained. The framework offers no account of why certain abilities appear discontinuously or why they require specific thresholds of model size. If emergent capabilities arise from implicit EM dynamics, the connection is not obvious; if they arise from other mechanisms, the framework is silent.

These are not failures of the analysis but boundaries of its scope. The contribution is to identify and derive one mechanism precisely, not to explain all of deep learning. Clarity about what is claimed protects against overinterpretation—and leaves room for complementary explanations of phenomena that implicit EM does not reach.

7 Discussion

The analysis is complete; what remains is to consider its significance. This section reflects on what the implicit EM framework unifies, what it implies for how we understand neural networks, and what questions it leaves open for future work.

7.1 Unification

The framework reveals that Gaussian mixture models, attention mechanisms, and cross-entropy classification are not three different methods with superficial similarities. They are one mechanism operating under different constraints. In GMMs, responsibilities are fully latent. In attention, responsibilities are conditioned on queries and recomputed per input. In cross-entropy, responsibilities are partially clamped by supervision. The underlying dynamics—exponentiation, normalization, responsibility-weighted updates—are identical.

This suggests a shift in how we interpret neural network training. Probability is often treated as primitive: we define distributions, derive likelihoods, and optimize. The implicit EM perspective inverts this. Distances are primary. Networks compute geometric quantities—deviations from learned structures—and probabilities arise only after exponentiation and normalization. Geometry precedes probability; inference is a consequence of optimization on geometric objectives.

Loss functions, in this view, are not arbitrary choices tuned for performance. They are geometric priors. Cross-entropy encodes an assumption that inputs should be assigned to discrete categories with full responsibility. Log-sum-exp over distances assumes inputs arise from a mixture of latent causes. Correntropy assumes outliers should be ignored. Each objective induces a different geometry of assignment and a different pattern of gradient flow. Choosing a loss function is choosing a theory of how data relates to structure—whether or not that choice is made consciously.

7.2 Implications

For interpretability, the framework offers a direct path from training dynamics to semantic structure. If responsibilities are gradients, then the assignments a network makes are not hidden quantities requiring probes or post-hoc analysis. They are present in the backward pass, computed at every training step. The question “which component is responsible for this input?” has an answer in the gradient itself. This does not solve interpretability—understanding *why* a component takes responsibility requires further analysis—but it locates the assignment structure in a quantity that is already computed, not one that must be extracted.

For objective design, the analysis reframes log-sum-exp not as a numerical convenience but as a structural requirement. Softmax is often introduced to avoid overflow or to produce well-behaved gradients. The implicit EM perspective reveals a deeper role: log-sum-exp is what induces competition, competition is what produces responsibilities, and responsibilities are what make learning behave like inference. If inference-like behavior is desired, LSE structure is not optional. If it is not desired—if independent predictions or robust outlier handling are preferred—then LSE should be deliberately avoided. The choice is not about numerical stability; it is about what kind of learning dynamics the objective will produce.

For theory, the framework dissolves a long-standing separation. Optimization and inference are traditionally distinct: one concerns finding parameters that minimize a loss, the other concerns computing posteriors over latent variables. The implicit EM result shows that under the right objectives, these are the same process viewed at different levels. Gradient descent *is* inference when the gradients *are* responsibilities. This is not a claim that all optimization is inference, but that for a well-defined class of objectives, the distinction disappears. The forward pass computes posteriors; the backward pass applies them. Training is inference, performed continuously over the dataset.

7.3 Open Directions

Several directions remain open. The absence of volume control in most neural objectives—the missing log-determinant—leads to collapse risks that are currently managed by heuristics. Section 3.4 supplies the exact term for the linear-layer case, where the volume correction is a log-determinant of the same weights that compute the distance. What remains open is the general case: deriving implicit volume terms for deep, nonlinear compositions of distance layers, and understanding when normalization layers substitute for explicit volume control and when they do not. Answering this would connect the implicit EM framework to practical stability concerns across architectures.

Supervision in real settings is rarely clean. Labels may be noisy, partial, or uncertain. The constrained regime analysis assumes hard labels that clamp responsibilities exactly; a fuller treatment

would model soft or probabilistic supervision as partial constraints on the responsibility structure. This could unify semi-supervised learning, label smoothing, and learning from crowds under the implicit EM framework.

Open-set inference requires escaping the closed-world assumption. Current objectives force every input to be assigned; realistic deployment requires the option to reject. Objectives that support non-assignment—an explicit “none of the above” component, or a threshold below which no component takes responsibility—would extend implicit EM to settings where not all inputs belong to known categories.

Finally, diagnostic tools are needed. If trained networks perform implicit EM, it should be possible to measure this: to extract responsibilities from gradients, to track specialization over training, to detect when the mechanism fails or degenerates. Such tools would move the framework from explanatory theory to practical instrument, enabling researchers to verify whether a given model exhibits the dynamics the theory predicts.

8 Conclusion

This paper establishes a single result: for objectives of the form $L = \log \sum_j \exp(-d_j)$, the gradient with respect to each distance is the negative responsibility of the corresponding component. The identity $\partial L / \partial d_j = -r_j$ is exact, requiring no approximations beyond differentiability.

The implication is immediate. Gradient descent on distance-based log-sum-exp objectives performs expectation-maximization implicitly. Responsibilities are not computed as auxiliary quantities; they are the gradients. The forward pass is the E-step; the backward pass is the M-step. No explicit inference algorithm is required because inference is already embedded in optimization.

This mechanism unifies phenomena that have been treated as distinct. Unsupervised mixture learning, attention in transformers, and cross-entropy classification are three regimes of the same underlying process—differing in what is observed and what is latent, but governed by identical dynamics. The Bayesian structure recently observed in trained transformers is not an emergent mystery; it is a necessary consequence of the objectives used to train them.

Optimization and inference are the same process viewed at different scales.

Acknowledgments

LLM assistants were used throughout this work as mathematical sounding boards and adversarial checkers—including the initial recognition that the log-sum-exp gradient reproduces EM responsibilities, and the identification of relevant classical literature. All derivations were verified independently, and the central identities are additionally verified numerically against finite differences in code accompanying this paper.

References

- Naman Agarwal, Siddhartha R. Dalal, and Vishal Misra. The Bayesian geometry of transformer attention. *arXiv preprint arXiv:2512.22471*, 2025a.
- Naman Agarwal, Siddhartha R. Dalal, and Vishal Misra. Gradient dynamics of attention: How cross-entropy sculpts Bayesian manifolds. *arXiv preprint arXiv:2512.22473*, 2025b.
- Naman Agarwal, Siddhartha R. Dalal, and Vishal Misra. Geometric scaling of Bayesian inference in LLMs. *arXiv preprint arXiv:2512.23752*, 2025c.

- Christopher M. Bishop. *Pattern Recognition and Machine Learning*. Springer, 2006.
- Olivier Cappé and Eric Moulines. On-line expectation–maximization algorithm for latent data models. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 71(3):593–613, 2009.
- Arthur P. Dempster, Nan M. Laird, and Donald B. Rubin. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society: Series B (Methodological)*, 39(1):1–38, 1977.
- Klaus Greff, Sjoerd van Steenkiste, and Jürgen Schmidhuber. Neural expectation maximization. In *Advances in Neural Information Processing Systems*, volume 30, pages 6694–6704, 2017.
- Geoffrey E. Hinton, Sara Sabour, and Nicholas Frosst. Matrix capsules with EM routing. In *International Conference on Learning Representations*, 2018.
- Robert A. Jacobs, Michael I. Jordan, Steven J. Nowlan, and Geoffrey E. Hinton. Adaptive mixtures of local experts. *Neural Computation*, 3(1):79–87, 1991.
- Yann LeCun, Sumit Chopra, Raia Hadsell, Marc’Aurelio Ranzato, and Fu Jie Huang. A tutorial on energy-based learning. In *Predicting Structured Data*, pages 191–246. MIT Press, 2006.
- Xia Li, Zhisheng Zhong, Jianlong Wu, Yibo Yang, Zhouchen Lin, and Hong Liu. Expectation-maximization attention networks for semantic segmentation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 9167–9176, 2019.
- Radford M. Neal and Geoffrey E. Hinton. A view of the EM algorithm that justifies incremental, sparse, and other variants. In Michael I. Jordan, editor, *Learning in Graphical Models*, pages 355–368. Springer, 1998.
- Alan Oursland. Neural networks use distance metrics. *arXiv preprint arXiv:2411.17932*, 2024a.
- Alan Oursland. Interpreting neural networks through Mahalanobis distance. *arXiv preprint arXiv:2410.19352*, 2024b.
- Alan Oursland. Neural networks learn distance metrics. *arXiv preprint arXiv:2502.02103*, 2025.
- Hubert Ramsauer, Bernhard Schäfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Thomas Adler, Lukas Gruber, Markus Holzleitner, Milena Pavlović, Geir Kjetil Sandve, Victor Greiff, David Kreil, Michael Kopp, Günter Klambauer, Johannes Brandstetter, and Sepp Hochreiter. Hopfield networks is all you need. In *International Conference on Learning Representations*, 2021.
- Danilo Jimenez Rezende and Shakir Mohamed. Variational inference with normalizing flows. In *International Conference on Machine Learning*, pages 1530–1538, 2015.
- Ruslan Salakhutdinov, Sam Roweis, and Zoubin Ghahramani. Optimization with EM and expectation-conjugate-gradient. In *International Conference on Machine Learning*, pages 672–679, 2003.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, pages 5998–6008, 2017.

- Lei Xu and Michael I. Jordan. On convergence properties of the EM algorithm for Gaussian mixtures. *Neural Computation*, 8(1):129–151, 1996.
- Guojun Zhang, Daniel Hsu, Arian Maleki, and Inderjit S. Dhillon. Comparing EM with GD in mixture models of two components. In *International Conference on Machine Learning*, pages 11285–11294. PMLR, 2020.